

Subsetting and assignment

Subsetting and assignment

Subsets can also be used with assignment to update specific values within an object.

```
x = c(1, 4, 7)
```

```
x[2] = 2  
x
```

```
## [1] 1 2 7
```

```
x %% 2 != 0
```

```
## [1] TRUE FALSE TRUE
```

```
x[x %% 2 != 0] = x[x %% 2 != 0] + 1  
x
```

```
## [1] 2 2 8
```

```
x[c(1,1)] = c(2,3)  
x
```

```
## [1] 3 2 8
```

```
x = 1:6
x[c(2,NA)] = 1
x
```

```
## [1] 1 1 3 4 5 6
```

```
x = 1:6
x[c(-1,-2)] = 3
x
```

```
## [1] 1 2 3 3 3 3
```

```
x = 1:6
x[c(TRUE,NA)] = 1
x
```

```
## [1] 1 2 1 4 1 6
```

```
x = 1:6
x[] = 1:3
x
```

```
## [1] 1 2 3 1 2 3
```



Tidy data

country	year	cases	population
Afghanistan	1999	17745	15557071
Afghanistan	2000	2666	20995360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	212258	1272915272
China	2000	210766	128042583

variables

country	year	cases	population
Afghanistan	1999	17745	15557071
Afghanistan	2000	2666	20995360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	212258	1272915272
China	2000	210766	128042583

observations

1. Every column is a variable.
2. Every row is an observation.
3. Every cell is a single value.

Tidy vs Untidy

Happy families are all alike; every unhappy family is unhappy in its own way


— Leo Tolstoy

```
tidyr::billboard[,1:7]
```

```
## # A tibble: 317 x 7
##   artist      track      date.entered  wk1  wk2  wk3  wk4
##   <chr>      <chr>      <date>      <dbl> <dbl> <dbl> <dbl>
## 1 2 Pac      Baby Don't Cry (Keep... 2000-02-26      87     82     72     77
## 2 2Ge+her    The Hardest Part Of ... 2000-09-02      91     87     92     NA
## 3 3 Doors Down Kryptonite      2000-04-08      81     70     68     67
## 4 3 Doors Down Loser           2000-10-21      76     76     72     69
## 5 504 Boyz   Wobble Wobble    2000-04-15      57     34     25     17
## 6 98^0      Give Me Just One Nig... 2000-08-19      51     39     34     26
## 7 A*Teens   Dancing Queen    2000-07-08      97     97     96     95
## 8 Aaliyah   I Don't Wanna    2000-01-29      84     62     51     41
## 9 Aaliyah   Try Again        2000-03-18      59     53     38     28
## 10 Adams, Yolanda Open My Heart    2000-08-26      76     76     74     69
## # ... with 307 more rows
```

Is the above data set tidy?

7133LE



www.rstudio.com

Modern data frames

Hadley Wickham / RStudio have a package that modifies data frames to be more modern, specifically they are designed to be surly and lazy + some fancy printing.

```
library(tibble)
```

```
iris
```

```
##      Sepal.Length Sepal.Width Petal.Length
## 1           5.1           3.5           1.4
## 2           4.9           3.0           1.4
## 3           4.7           3.2           1.3
## 4           4.6           3.1           1.5
## 5           5.0           3.6           1.4
## 6           5.4           3.9           1.7
## 7           4.6           3.4           1.4
## 8           5.0           3.4           1.5
## 9           4.4           2.9           1.4
## 10          4.9           3.1           1.5
## 11          5.4           3.7           1.5
## 12          4.8           3.4           1.6
## 13          4.8           3.0           1.4
## 14          4.3           3.0           1.1
## 15          5.8           4.0           1.2
## 16          5.7           4.4           1.5
## 17          5.4           3.9           1.3
## 18          5.1           3.5           1.4
```

```
(tbl_iris = as_tibble(iris))
```

```
## # A tibble: 150 x 5
##   Sepal.Length Sepal.Width Petal.Length
##         <dbl>         <dbl>         <dbl>
## 1           5.1           3.5           1.4
## 2           4.9           3           1.4
## 3           4.7           3.2           1.3
## 4           4.6           3.1           1.5
## 5           5           3.6           1.4
## 6           5.4           3.9           1.7
## 7           4.6           3.4           1.4
## 8           5           3.4           1.5
## 9           4.4           2.9           1.4
## 10          4.9           3.1           1.5
## # ... with 140 more rows, and 2 more variables:
## #   Petal.Width <dbl>, Species <fct>
```


Tibbles are lazy

By default, tibbles will always stay as tibbles when subsetting (except when using `$` or `[]`).

```
tbl_iris[1,]
```

```
## # A tibble: 1 x 5
##   Sepal.Length Sepal.Width Petal.Length Petal.Width Species
##   <dbl>         <dbl>         <dbl>         <dbl> <fct>
## 1         5.1         3.5         1.4         0.2 setosa
```

```
tbl_iris[,1]
```

```
## # A tibble: 150 x 1
##   Sepal.Length
##   <dbl>
## 1         5.1
## 2         4.9
## 3         4.7
## 4         4.6
## 5          5
## 6         5.4
## 7         4.6
## 8          5
## 9         4.4
## 10        4.9
## # ... with 140 more rows
```

```
tbl_iris[[1]]
```

```
##   [1] 5.1 4.9 4.7 4.6 5.0 5.4 4.6 5.0 4.4 4.9 5.4
##  [19] 5.7 5.1 5.4 5.1 4.6 5.1 4.8 5.0 5.0 5.2 5.2
##  [37] 5.5 4.9 4.4 5.1 5.0 4.5 4.4 5.0 5.1 4.8 5.1
##  [55] 6.5 5.7 6.3 4.9 6.6 5.2 5.0 5.9 6.0 6.1 5.6
##  [73] 6.3 6.1 6.4 6.6 6.8 6.7 6.0 5.7 5.5 5.5 5.8
##  [91] 5.5 6.1 5.8 5.0 5.6 5.7 5.7 6.2 5.1 5.7 6.3
## [109] 6.7 7.2 6.5 6.4 6.8 5.7 5.8 6.4 6.5 7.7 7.7
## [127] 6.2 6.1 6.4 7.2 7.4 7.9 6.4 6.3 6.1 7.7 6.3
## [145] 6.7 6.7 6.3 6.5 6.2 5.9
```

More laziness - partial matching

Tibbles do not use partial matching when the `$` operator is used.

```
head( iris$Sp )
```

```
## [1] setosa setosa setosa setosa setosa setosa  
## Levels: setosa versicolor virginica
```

```
head( iris$Species )
```

```
## [1] setosa setosa setosa setosa setosa setosa  
## Levels: setosa versicolor virginica
```

```
head( tbl_iris$Sp )
```

```
## Warning: Unknown or uninitialised column: `Sp`.  
## NULL
```

```
head( tbl_iris$Species )
```

```
## [1] setosa setosa setosa setosa setosa setosa  
## Levels: setosa versicolor virginica
```

More laziness - stringsAsFactors

Tibbles also have always had `stringsAsFactors = FALSE` as default behavior.

```
(t = tibble(  
  x = 1:3,  
  y = c("A", "B", "C"),  
  z = factor(c("X", "Y", "Z"))  
))
```

```
## # A tibble: 3 x 3  
##       x y     z  
##   <int> <chr> <fct>  
## 1     1 A     X  
## 2     2 B     Y  
## 3     3 C     Z
```

```
str(t)
```

```
## tibble [3 × 3] (S3: tbl_df/tbl/data.frame)  
## $ x: int [1:3] 1 2 3  
## $ y: chr [1:3] "A" "B" "C"  
## $ z: Factor w/ 3 levels "X","Y","Z": 1 2 3
```

```
(d = data.frame(  
  x = 1:3,  
  y = c("A", "B", "C"),  
  z = factor(c("X", "Y", "Z")),  
  stringsAsFactors = TRUE  
))
```

```
##      x y z  
## 1 1 A X  
## 2 2 B Y  
## 3 3 C Z
```

```
str(d)
```

```
## 'data.frame':    3 obs. of  3 variables:  
## $ x: int  1 2 3  
## $ y: Factor w/ 3 levels "A","B","C": 1 2 3  
## $ z: Factor w/ 3 levels "X","Y","Z": 1 2 3
```

Tibbles and length coercion

```
tibble(x = 1:4, y = 1)
```

```
## # A tibble: 4 x 2
##       x     y
##   <int> <dbl>
## 1     1     1
## 2     2     1
## 3     3     1
## 4     4     1
```

```
tibble(x = 1:4, y = 1:2)
```

```
## Error: Tibble columns must have compatible sizes.
## * Size 4: Existing data.
## * Size 2: Column `y`.
## i Only values of size one are recycled.
```

```
tibble(x = 1:4, y = 1:3)
```

```
## Error: Tibble columns must have compatible sizes.
## * Size 4: Existing data.
## * Size 3: Column `y`.
## i Only values of size one are recycled.
```



magrittr

Pipes in R

You can think about the following sequence of actions - find key, unlock car, start car, drive to school, park.

Expressed as a set of nested functions in R pseudocode this would look like:

```
park(drive(start_car(find("keys")), to="campus"))
```

Writing it out using pipes give it a more natural (and easier to read) structure:

```
find("keys") %>%  
  start_car() %>%  
  drive(to="campus") %>%  
  park()
```

Approaches

All of the following are fine, it comes down to personal preference:

Nested:

```
h( g( f(x), y=1), z=1 )
```

Piped:

```
f(x) %>% g(y=1) %>% h(z=1)
```

Intermediate:

```
res = f(x)  
res = g(res, y=1)  
res = h(res, z=1)
```

What about other arguments?

Sometimes we want to send our results to an function argument other than first one or we want to use the previous result for multiple arguments. In these cases we can refer to the previous result using `..`

```
data.frame(a = 1:3, b = 3:1) %>% lm(a~b, data=.)
```

```
##  
## Call:  
## lm(formula = a ~ b, data = .)  
##  
## Coefficients:  
## (Intercept)          b  
##           4          -1
```

```
data.frame(a = 1:3, b = 3:1) %>% .[[1]]
```

```
## [1] 1 2 3
```

```
data.frame(a = 1:3, b = 3:1) %>% .[[length(.)]]
```

```
## [1] 3 2 1
```


A Grammar of Data Manipulation

dplyr is based on the concepts of functions as verbs that manipulate data frames.

Single data frame functions / verbs:

- `filter()` / `slice()`: pick rows based on criteria
- `select()` / `rename()`: select columns by name
- `pull()`: grab a column as a vector
- `arrange()`: reorder rows
- `mutate()` / `transmute()`: create or modify columns
- `distinct()`: filter for unique rows
- `summarise()` / `count()`: reduce variables to values
- `group_by()` / `ungroup()`: modify other verbs to act on subsets
- `relocate()`: change column order
- ... (many more)

dplyr rules

1. First argument is *always* a data frame
2. Subsequent arguments say what to do with that data frame
3. *Always* return a data frame
4. Don't modify in place
5. Lazy evaluation magic

Example Data

We will demonstrate dplyr's functionality using the nycflights13 data.

```
library(dplyr)
library(nycflights13)
```

```
flights
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>     <int>         <int>
## 1  2013     1     1     517           515           2       830           819
## 2  2013     1     1     533           529           4       850           830
## 3  2013     1     1     542           540           2       923           850
## 4  2013     1     1     544           545          -1      1004          1022
## 5  2013     1     1     554           600          -6       812           837
## 6  2013     1     1     554           558          -4       740           728
## 7  2013     1     1     555           600          -5       913           854
## 8  2013     1     1     557           600          -3       709           723
## 9  2013     1     1     557           600          -3       838           846
## 10 2013     1     1     558           600          -2       753           745
## # ... with 336,766 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

filter() - March flights

```
flights %>% filter(month == 3)
```

```
## # A tibble: 28,834 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>     <int>         <int>
## 1  2013     3     1         4           2159          125       318           56
## 2  2013     3     1        50           2358           52       526          438
## 3  2013     3     1       117           2245          152       223         2354
## 4  2013     3     1       454           500           -6       633          648
## 5  2013     3     1      505           515          -10       746          810
## 6  2013     3     1      521           530           -9       813          827
## 7  2013     3     1      537           540            -3       856          850
## 8  2013     3     1      541           545            -4      1014         1023
## 9  2013     3     1      549           600          -11       639          703
## 10 2013     3     1      550           600          -10       747          801
## # ... with 28,824 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

filter() - Flights in the first 7 days of March

```
flights %>% filter(month == 3, day <= 7)
```

```
## # A tibble: 6,530 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     3     1         4           2159          125     318           56
## 2  2013     3     1        50           2358           52     526          438
## 3  2013     3     1       117           2245          152     223         2354
## 4  2013     3     1       454           500           -6     633           648
## 5  2013     3     1      505           515          -10     746           810
## 6  2013     3     1      521           530           -9     813           827
## 7  2013     3     1      537           540           -3     856           850
## 8  2013     3     1      541           545           -4    1014          1023
## 9  2013     3     1      549           600          -11     639           703
## 10 2013     3     1      550           600          -10     747           801
## # ... with 6,520 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

filter() - Flights to LAX or JFK in March

```
flights %>% filter(dest == "LAX" | dest == "JFK", month==3)
```

```
## # A tibble: 1,178 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>    <int>         <int>
## 1  2013     3     1     607           610          -3      832           925
## 2  2013     3     1     629           632          -3      844           952
## 3  2013     3     1     657           700          -3      953          1034
## 4  2013     3     1     714           715          -1      939          1037
## 5  2013     3     1     716           710           6      958          1035
## 6  2013     3     1     727           730          -3     1007          1100
## 7  2013     3     1     836           840          -4     1111          1157
## 8  2013     3     1     857           900          -3     1202          1221
## 9  2013     3     1     903           900           3     1157          1220
## 10 2013     3     1     904           831          33     1150          1151
## # ... with 1,168 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

slice() - First 10 flights

```
flights %>% slice(1:10)
```

```
## # A tibble: 10 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517           515           2     830           819
## 2  2013     1     1     533           529           4     850           830
## 3  2013     1     1     542           540           2     923           850
## 4  2013     1     1     544           545          -1    1004          1022
## 5  2013     1     1     554           600          -6     812           837
## 6  2013     1     1     554           558          -4     740           728
## 7  2013     1     1     555           600          -5     913           854
## 8  2013     1     1     557           600          -3     709           723
## 9  2013     1     1     557           600          -3     838           846
## 10 2013     1     1     558           600          -2     753           745
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```


slice() - Last 5 flights

```
flights %>% slice((n()-4):n())
```

```
## # A tibble: 5 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     9    30      NA           1455          NA      NA           1634
## 2  2013     9    30      NA           2200          NA      NA           2312
## 3  2013     9    30      NA           1210          NA      NA           1330
## 4  2013     9    30      NA           1159          NA      NA           1344
## 5  2013     9    30      NA            840          NA      NA           1020
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
flights %>% slice_tail(n = 5)
```

```
## # A tibble: 5 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     9    30      NA           1455          NA      NA           1634
## 2  2013     9    30      NA           2200          NA      NA           2312
## 3  2013     9    30      NA           1210          NA      NA           1330
## 4  2013     9    30      NA           1159          NA      NA           1344
## 5  2013     9    30      NA            840          NA      NA           1020
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

select() - Individual Columns

```
flights %>% select(year, month, day)
```

```
## # A tibble: 336,776 x 3
##   year month   day
##   <int> <int> <int>
## 1  2013     1     1
## 2  2013     1     1
## 3  2013     1     1
## 4  2013     1     1
## 5  2013     1     1
## 6  2013     1     1
## 7  2013     1     1
## 8  2013     1     1
## 9  2013     1     1
## 10 2013     1     1
## # ... with 336,766 more rows
```

select() - Exclude Columns

```
flights %>% select(-year, -month, -day)
```

```
## # A tibble: 336,776 x 16
##   dep_time sched_dep_time dep_delay arr_time sched_arr_time arr_delay carrier
##   <int>      <int>      <dbl>   <int>      <int>      <dbl>   <chr>
## 1      517         515         2      830         819        11   UA
## 2      533         529         4      850         830        20   UA
## 3      542         540         2      923         850        33   AA
## 4      544         545        -1     1004        1022       -18   B6
## 5      554         600        -6      812         837       -25   DL
## 6      554         558        -4      740         728        12   UA
## 7      555         600        -5      913         854        19   B6
## 8      557         600        -3      709         723       -14   EV
## 9      557         600        -3      838         846        -8   B6
## 10     558         600        -2      753         745         8   AA
## # ... with 336,766 more rows, and 9 more variables: flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

select() - Ranges

```
flights %>% select(year:day)
```

```
## # A tibble: 336,776 x 3
##   year month   day
##   <int> <int> <int>
## 1  2013     1     1
## 2  2013     1     1
## 3  2013     1     1
## 4  2013     1     1
## 5  2013     1     1
## 6  2013     1     1
## 7  2013     1     1
## 8  2013     1     1
## 9  2013     1     1
## 10 2013     1     1
## # ... with 336,766 more rows
```

select() - Exclusion Ranges

```
flights %>% select(-(year:day))
```

```
## # A tibble: 336,776 x 16
##   dep_time sched_dep_time dep_delay arr_time sched_arr_time arr_delay carrier
##   <int>      <int>      <dbl>    <int>      <int>      <dbl> <chr>
## 1      517         515         2      830         819        11 UA
## 2      533         529         4      850         830        20 UA
## 3      542         540         2      923         850        33 AA
## 4      544         545        -1     1004        1022       -18 B6
## 5      554         600        -6      812         837       -25 DL
## 6      554         558        -4      740         728        12 UA
## 7      555         600        -5      913         854        19 B6
## 8      557         600        -3      709         723       -14 EV
## 9      557         600        -3      838         846        -8 B6
## 10     558         600        -2      753         745         8 AA
## # ... with 336,766 more rows, and 9 more variables: flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

select() - Matching

```
flights %>% select(contains("dep"),
                  contains("arr"))
```

```
## # A tibble: 336,776 x 7
##   dep_time sched_dep_time dep_delay arr_time sched_arr_time arr_delay carrier
##   <int>         <int>         <dbl>   <int>         <int>         <dbl>   <chr>
## 1      517           515           2      830           819          11    UA
## 2      533           529           4      850           830          20    UA
## 3      542           540           2      923           850          33    AA
## 4      544           545          -1     1004          1022         -18    B6
## 5      554           600          -6      812           837         -25    DL
## 6      554           558          -4      740           728          12    UA
## 7      555           600          -5      913           854          19    B6
## 8      557           600          -3      709           723         -14    EV
## 9      557           600          -3      838           846           -8    B6
## 10     558           600          -2      753           745           8    AA
## # ... with 336,766 more rows
```

```
flights %>% select(starts_with("dep"),
                  starts_with("arr"))
```

```
## # A tibble: 336,776 x 4
##   dep_time dep_delay arr_time arr_delay
##   <int>     <dbl>    <int>    <dbl>
## 1     517         2      830        11
## 2     533         4      850        20
## 3     542         2      923        33
## 4     544        -1     1004       -18
## 5     554        -6      812       -25
## 6     554        -4      740        12
## 7     555        -5      913        19
## 8     557        -3      709       -14
## 9     557        -3      838        -8
## 10    558        -2      753         8
## # ... with 336,766 more rows
```

Other helpers provide by tidyselect: `starts_with`, `ends_with`, `everything`, `matches`, `num_range`, `one_of`, `everything`, `last_col`.

select_if() - Get non-numeric columns

```
flights %>% select_if(function(x) !is.numeric(x))
```

```
## # A tibble: 336,776 x 5
##   carrier tailnum origin dest time_hour
##   <chr>    <chr>   <chr> <chr> <dtm>
## 1 UA      N14228   EWR   IAH   2013-01-01 05:00:00
## 2 UA      N24211   LGA   IAH   2013-01-01 05:00:00
## 3 AA      N619AA    JFK   MIA   2013-01-01 05:00:00
## 4 B6      N804JB    JFK   BQN   2013-01-01 05:00:00
## 5 DL      N668DN    LGA   ATL   2013-01-01 06:00:00
## 6 UA      N39463    EWR   ORD   2013-01-01 05:00:00
## 7 B6      N516JB    EWR   FLL   2013-01-01 06:00:00
## 8 EV      N829AS    LGA   IAD   2013-01-01 06:00:00
## 9 B6      N593JB    JFK   MCO   2013-01-01 06:00:00
## 10 AA     N3ALAA    LGA   ORD   2013-01-01 06:00:00
## # ... with 336,766 more rows
```

relocate - to the front

```
flights %>% relocate(carrier, origin, dest)
```

```
## # A tibble: 336,776 x 19
##   carrier origin dest   year month   day dep_time sched_dep_time dep_delay
##   <chr>    <chr> <chr> <int> <int> <int>   <int>         <int>         <dbl>
## 1 UA      EWR    IAH   2013     1     1     517           515             2
## 2 UA      LGA    IAH   2013     1     1     533           529             4
## 3 AA      JFK    MIA   2013     1     1     542           540             2
## 4 B6      JFK    BQN   2013     1     1     544           545            -1
## 5 DL      LGA    ATL   2013     1     1     554           600            -6
## 6 UA      EWR    ORD   2013     1     1     554           558            -4
## 7 B6      EWR    FLL   2013     1     1     555           600            -5
## 8 EV      LGA    IAD   2013     1     1     557           600            -3
## 9 B6      JFK    MCO   2013     1     1     557           600            -3
## 10 AA     LGA    ORD   2013     1     1     558           600            -2
## # ... with 336,766 more rows, and 10 more variables: arr_time <int>,
## #   sched_arr_time <int>, arr_delay <dbl>, flight <int>, tailnum <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

relocate - to the end

```
flights %>%  
  relocate(year, month, day, .after = last_col())
```

```
## # A tibble: 336,776 x 19  
##   dep_time sched_dep_time dep_delay arr_time sched_arr_time arr_delay carrier  
##   <int>         <int>         <dbl>   <int>         <int>         <dbl> <chr>  
## 1      517           515           2     830           819          11 UA  
## 2      533           529           4     850           830          20 UA  
## 3      542           540           2     923           850          33 AA  
## 4      544           545          -1    1004          1022         -18 B6  
## 5      554           600          -6     812           837         -25 DL  
## 6      554           558          -4     740           728          12 UA  
## 7      555           600          -5     913           854          19 B6  
## 8      557           600          -3     709           723         -14 EV  
## 9      557           600          -3     838           846          -8 B6  
## 10     558           600          -2     753           745           8 AA  
## # ... with 336,766 more rows, and 12 more variables: flight <int>, tailnum <chr>,  
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,  
## #   minute <dbl>, time_hour <dtm>, year <int>, month <int>, day <int>
```

rename() - Change column names

```
flights %>% rename(tail_number = tailnum)
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517           515           2     830           819
## 2  2013     1     1     533           529           4     850           830
## 3  2013     1     1     542           540           2     923           850
## 4  2013     1     1     544           545          -1    1004          1022
## 5  2013     1     1     554           600          -6     812           837
## 6  2013     1     1     554           558          -4     740           728
## 7  2013     1     1     555           600          -5     913           854
## 8  2013     1     1     557           600          -3     709           723
## 9  2013     1     1     557           600          -3     838           846
## 10 2013     1     1     558           600          -2     753           745
## # ... with 336,766 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tail_number <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

select() vs. rename()

```
flights %>% select(tail_number = tailnum)
```

```
## # A tibble: 336,776 x 1
##   tail_number
##   <chr>
## 1 N14228
## 2 N24211
## 3 N619AA
## 4 N804JB
## 5 N668DN
## 6 N39463
## 7 N516JB
## 8 N829AS
## 9 N593JB
## 10 N3ALAA
## # ... with 336,766 more rows
```

```
flights %>% rename(tail_number = tailnum)
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517             515           2     830             819
## 2  2013     1     1     533             529           4     850             830
## 3  2013     1     1     542             540           2     923             850
## 4  2013     1     1     544             545          -1    1004            1022
## 5  2013     1     1     554             600          -6     812             837
## 6  2013     1     1     554             558          -4     740             728
```

pull()

```
names(flights)
```

```
## [1] "year"      "month"      "day"        "dep_time"  
## [5] "sched_dep_time" "dep_delay"  "arr_time"   "sched_arr_time"  
## [9] "arr_delay"    "carrier"    "flight"     "tailnum"  
## [13] "origin"      "dest"       "air_time"   "distance"  
## [17] "hour"       "minute"     "time_hour"
```

```
flights %>% pull("year") %>% head()
```

```
## [1] 2013 2013 2013 2013 2013 2013
```

```
flights %>% pull(1) %>% head()
```

```
## [1] 2013 2013 2013 2013 2013 2013
```

```
flights %>% pull(-1) %>% head()
```

```
## [1] "2013-01-01 05:00:00 EST" "2013-01-01 05:00:00 EST"  
## [3] "2013-01-01 05:00:00 EST" "2013-01-01 05:00:00 EST"  
## [5] "2013-01-01 06:00:00 EST" "2013-01-01 05:00:00 EST"
```

arrange() - Sort data

```
flights %>% filter(month==3,day==2) %>% arrange(origin, dest)
```

```
## # A tibble: 765 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     3     2    1336         1329           7    1426         1432
## 2  2013     3     2     628           629          -1     837           849
## 3  2013     3     2     637           640          -3     903           915
## 4  2013     3     2     743           745          -2     945          1010
## 5  2013     3     2     857           900          -3    1117          1126
## 6  2013     3     2    1027          1030          -3    1234          1247
## 7  2013     3     2    1134          1145         -11    1332          1359
## 8  2013     3     2    1412          1415          -3    1636          1630
## 9  2013     3     2    1633          1636          -3    1848          1908
## 10 2013     3     2    1655          1700          -5    1857          1924
## # ... with 755 more rows, and 11 more variables: arr_delay <dbl>, carrier <chr>,
## #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>,
## #   distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

arrange() & desc() - Descending order

```
flights %>%  
  filter(month==3, day==2) %>%  
  arrange(desc(origin), dest) %>%  
  select(origin, dest, tailnum)
```

```
## # A tibble: 765 x 3  
##   origin dest tailnum  
##   <chr> <chr> <chr>  
## 1 LGA    ATL   N928AT  
## 2 LGA    ATL   N623DL  
## 3 LGA    ATL   N680DA  
## 4 LGA    ATL   N996AT  
## 5 LGA    ATL   N510MQ  
## 6 LGA    ATL   N663DN  
## 7 LGA    ATL   N942DL  
## 8 LGA    ATL   N511MQ  
## 9 LGA    ATL   N910DE  
## 10 LGA   ATL   N902DE  
## # ... with 755 more rows
```


mutate() - Modify columns

```
flights %>%  
  select(year:day) %>%  
  mutate(date = paste(year, month, day, sep="/"))
```

```
## # A tibble: 336,776 x 4  
##   year month   day date  
##   <int> <int> <int> <chr>  
## 1  2013     1     1 2013/1/1  
## 2  2013     1     1 2013/1/1  
## 3  2013     1     1 2013/1/1  
## 4  2013     1     1 2013/1/1  
## 5  2013     1     1 2013/1/1  
## 6  2013     1     1 2013/1/1  
## 7  2013     1     1 2013/1/1  
## 8  2013     1     1 2013/1/1  
## 9  2013     1     1 2013/1/1  
## 10 2013     1     1 2013/1/1  
## # ... with 336,766 more rows
```

distinct() - Find unique rows

```
flights %>%  
  select(origin, dest) %>%  
  distinct() %>%  
  arrange(origin,dest)
```

```
## # A tibble: 224 x 2  
##   origin dest  
##   <chr>  <chr>  
## 1 EWR    ALB  
## 2 EWR    ANC  
## 3 EWR    ATL  
## 4 EWR    AUS  
## 5 EWR    AVL  
## 6 EWR    BDL  
## 7 EWR    BNA  
## 8 EWR    BOS  
## 9 EWR    BQN  
## 10 EWR   BTV  
## # ... with 214 more rows
```

summarise()

```
flights %>%  
  summarize(n(), min(dep_delay), max(dep_delay))
```

```
## # A tibble: 1 x 3  
##   `n()` `min(dep_delay)` `max(dep_delay)`  
##   <int>      <dbl>      <dbl>  
## 1 336776         NA         NA
```

```
flights %>%  
  summarize(  
    n = n(),  
    min_dep_delay = min(dep_delay, na.rm = TRUE),  
    max_dep_delay = max(dep_delay, na.rm = TRUE)  
  )
```

```
## # A tibble: 1 x 3  
##       n min_dep_delay max_dep_delay  
##   <int>      <dbl>      <dbl>  
## 1 336776        -43        1301
```

group_by()

```
flights %>% group_by(origin)
```

```
## # A tibble: 336,776 x 19
## # Groups:   origin [3]
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517           515           2     830           819
## 2  2013     1     1     533           529           4     850           830
## 3  2013     1     1     542           540           2     923           850
## 4  2013     1     1     544           545          -1    1004          1022
## 5  2013     1     1     554           600          -6     812           837
## 6  2013     1     1     554           558          -4     740           728
## 7  2013     1     1     555           600          -5     913           854
## 8  2013     1     1     557           600          -3     709           723
## 9  2013     1     1     557           600          -3     838           846
## 10 2013     1     1     558           600          -2     753           745
## # ... with 336,766 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

summarise() with group_by()

```
flights %>%  
  group_by(origin) %>%  
  summarize(  
    n = n(),  
    min_dep_delay = min(dep_delay, na.rm = TRUE),  
    max_dep_delay = max(dep_delay, na.rm = TRUE)  
  )
```

`summarise()` ungrouping output (override with `.groups` argument)

```
## # A tibble: 3 x 4  
##   origin      n min_dep_delay max_dep_delay  
##   <chr>   <int>         <dbl>         <dbl>  
## 1 EWR    120835          -25           1126  
## 2 JFK    111279          -43           1301  
## 3 LGA    104662          -33            911
```

summarise() with group_by() - quietly

```
flights %>%
  group_by(origin) %>%
  summarize(
    n = n(),
    min_dep_delay = min(dep_delay, na.rm = TRUE),
    max_dep_delay = max(dep_delay, na.rm = TRUE),
    .groups = "drop_last"
  )
```

```
## # A tibble: 3 x 4
##   origin      n min_dep_delay max_dep_delay
##   <chr>    <int>         <dbl>         <dbl>
## 1 EWR     120835          -25           1126
## 2 JFK     111279          -43           1301
## 3 LGA     104662          -33            911
```

```
flights %>%
  group_by(origin) %>%
  summarize(
    n = n(),
    min_dep_delay = min(dep_delay, na.rm = TRUE),
    max_dep_delay = max(dep_delay, na.rm = TRUE),
    .groups = "keep"
  )
```

```
## # A tibble: 3 x 4
## # Groups:   origin [3]
##   origin      n min_dep_delay max_dep_delay
##   <chr>    <int>         <dbl>         <dbl>
## 1 EWR     120835          -25           1126
## 2 JFK     111279          -43           1301
## 3 LGA     104662          -33            911
```

```

flights %>%
  group_by(origin, carrier) %>%
  summarize(
    n = n(),
    min_dep_delay = min(dep_delay, na.rm = TRUE),
    max_dep_delay = max(dep_delay, na.rm = TRUE)
  ) %>%
  filter(n > 10000)

```

```
## `summarise()` regrouping output by 'origin' (override with `.groups` argument)
```

```
## # A tibble: 10 x 5
```

```
## # Groups:   origin [3]
```

	origin	carrier	n	min_dep_delay	max_dep_delay
	<chr>	<chr>	<int>	<dbl>	<dbl>
## 1	EWB	EV	43939	-25	548
## 2	EWB	UA	46087	-18	424
## 3	JFK	9E	14651	-24	747
## 4	JFK	AA	13783	-15	1014
## 5	JFK	B6	42076	-43	453
## 6	JFK	DL	20701	-18	960
## 7	LGA	AA	15459	-24	803
## 8	LGA	DL	23067	-33	911
## 9	LGA	MQ	16928	-26	366
## 10	LGA	US	13136	-18	500

count()

```
flights %>%  
  group_by(origin, carrier) %>%  
  summarize(n = n(), .groups = "drop")
```

```
## # A tibble: 35 x 3  
##   origin carrier      n  
##   <chr>   <chr>   <int>  
## 1 EWR     9E     1268  
## 2 EWR     AA     3487  
## 3 EWR     AS      714  
## 4 EWR     B6     6557  
## 5 EWR     DL     4342  
## 6 EWR     EV    43939  
## 7 EWR     MQ     2276  
## 8 EWR     OO        6  
## 9 EWR     UA    46087  
## 10 EWR    US     4405  
## # ... with 25 more rows
```

```
flights %>%  
  count(origin, carrier)
```

```
## # A tibble: 35 x 3  
##   origin carrier      n  
##   <chr>   <chr>   <int>  
## 1 EWR     9E     1268  
## 2 EWR     AA     3487  
## 3 EWR     AS      714  
## 4 EWR     B6     6557  
## 5 EWR     DL     4342  
## 6 EWR     EV    43939  
## 7 EWR     MQ     2276  
## 8 EWR     OO        6  
## 9 EWR     UA    46087  
## 10 EWR    US     4405  
## # ... with 25 more rows
```


mutate() with group_by()

```
flights %>% group_by(origin) %>%  
  mutate(  
    n = n(),  
  ) %>%  
  select(origin, n)
```

```
## # A tibble: 336,776 x 2  
## # Groups:   origin [3]  
##   origin      n  
##   <chr>   <int>  
## 1 EWR    120835  
## 2 LGA    104662  
## 3 JFK    111279  
## 4 JFK    111279  
## 5 LGA    104662  
## 6 EWR    120835  
## 7 EWR    120835  
## 8 LGA    104662  
## 9 JFK    111279  
## 10 LGA    104662  
## # ... with 336,766 more rows
```

Demos

1. How many flights to Los Angeles (LAX) did each of the legacy carriers (AA, UA, DL or US) have in May from JFK, and what was their average duration?
2. What was the shortest flight out of each airport in terms of distance? In terms of duration?

Exercise 1

1. Which plane (check the tail number) flew out of each New York airport the most?
2. Which date should you fly on if you want to have the lowest possible average departure delay? What about arrival delay?